

# Today's topics

- Introductions
- Learning goals and objectives
- Course syllabus
- Exploring the classroom/C programming

# I bit about myself

- My name is Taha Khan
  - It's my fourth year at Washington and Lee University (W&L)
  - You may call me Taha or Professor/Dr. Khan
- My research focuses on how can we make the Internet more secure, robust, and usable.
- How to contact me?
  - Office: Parmly 406
  - Slack for Online Communications
  - Office Hours: M/W/F from 2:30 pm – 3:30 pm

# Class introductions

- What is your preferred name?
- What do you plan to get out of this course?
- What is the lowest level computer language you've programmed in?
- Tell us all something fun you did over the summers?

# Systems programming

- What is a computer system?
- Themes of this course:
  - End to end process of compiling and executing programs
  - Inspect running programs and low-level debugging
  - Memory management and virtual memory
  - The details of I/O
  - Coordinating with other processes

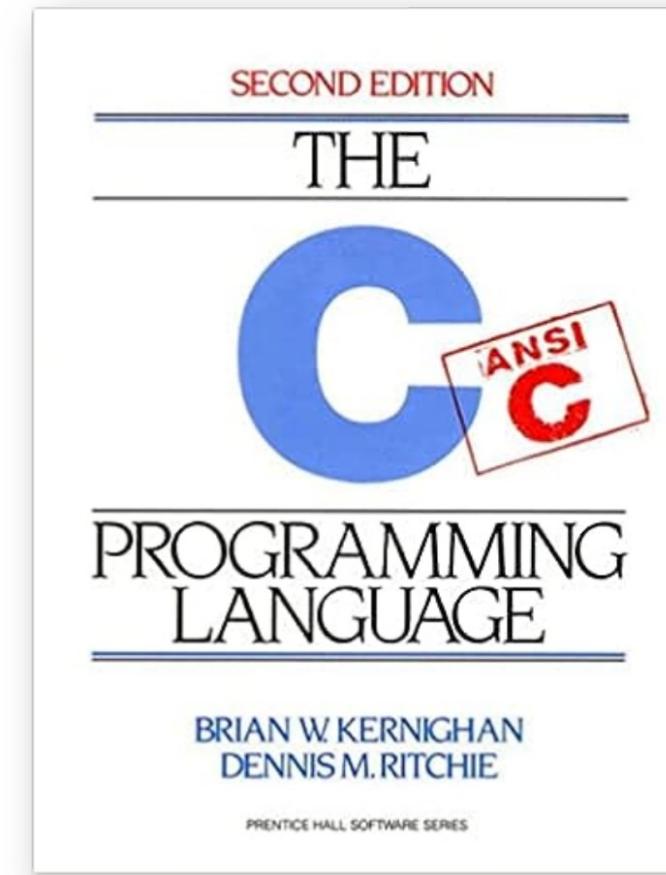
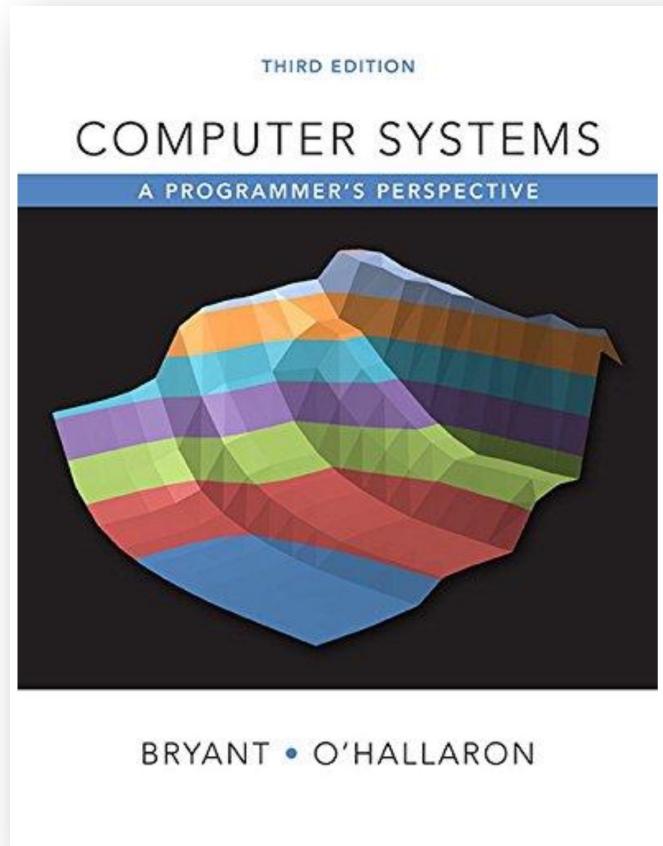
# Why C programming?

- C is the primary language that is used today for UNIX based systems
  - Standardization and portability
  - Low level capabilities, yet a human readable syntax
  - Generates very efficient code
- It has stood the test of time
- Roughly 80% of the computers out there are running UNIX systems
- Other benefits
  - Opens doors to many other avenues in computer science
  - A highly desired skill in the CS job market

# Course components and grading

- Homework assignments (35%)
  - Midterm exam (20%)
  - Final exam (30%)
  - Class presentations (10%)
  - In class activities (5%)
  - Bonus points (up to 2%)
- 
- Grades will be assigned on a relative curve
  - You have a quota of 5 late days that you may use for homework assignments

# Textbooks

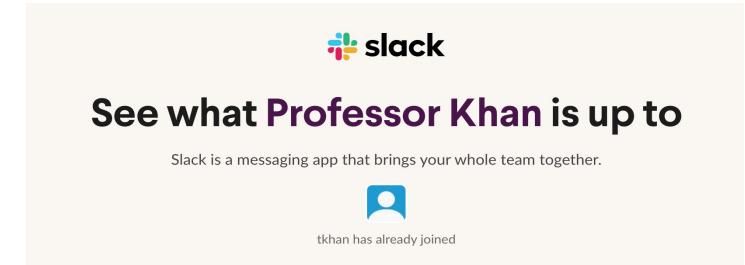


# Academic integrity

- Collaboration
  - You may discuss approaches, not solutions
  - You must submit your own work
  - Exams may include questions on programming
- Cheating
  - Zero tolerance for cheating, don't do it!
  - Potential violations of the W&L Honor System will be reported to the EC
  - All cheating results in AT LEAST one letter grade reduction
- You may use generative AI for studying concepts, but not to generate code

# Resources and communications

- W&L Canvas
  - Course materials
  - Homework assignment submissions
  - Grades
- Gradescope
  - Exams
- Slack
  - Communications
  - Please use your **Full Name in Title Case** when signing up



We suggest using the email account you use for work.

Email

name@work-email.com

Full name

Your name

Continue with Email

It's okay to send me emails about Slack.

By continuing, you're agreeing to our [User Terms of Service](#), [Privacy Policy](#), and [Cookie Policy](#).

# Bring your laptops to class

- In class programming exercises
- Exploration of tools
- AWS for homework assignments

# C programming exercises

1. Write a simple C function that swaps two integer variables using pointers
  - Test your function in `main()`
  - Compile and run your program using `gcc`
2. Next, modify your function (above), so that it is able to swap variables of any data type
  - Make sure to test it as well
3. Write a function that converts Celsius temperature to Fahrenheit
  - The program should prompt the user for input temperature in Celsius
  - You may compute everything in the main function or write a helper

# Next Class

- **Topic:** C programming basics
- **Reading:** Chapter 1, 2 and 3 of the C Book (K&R)

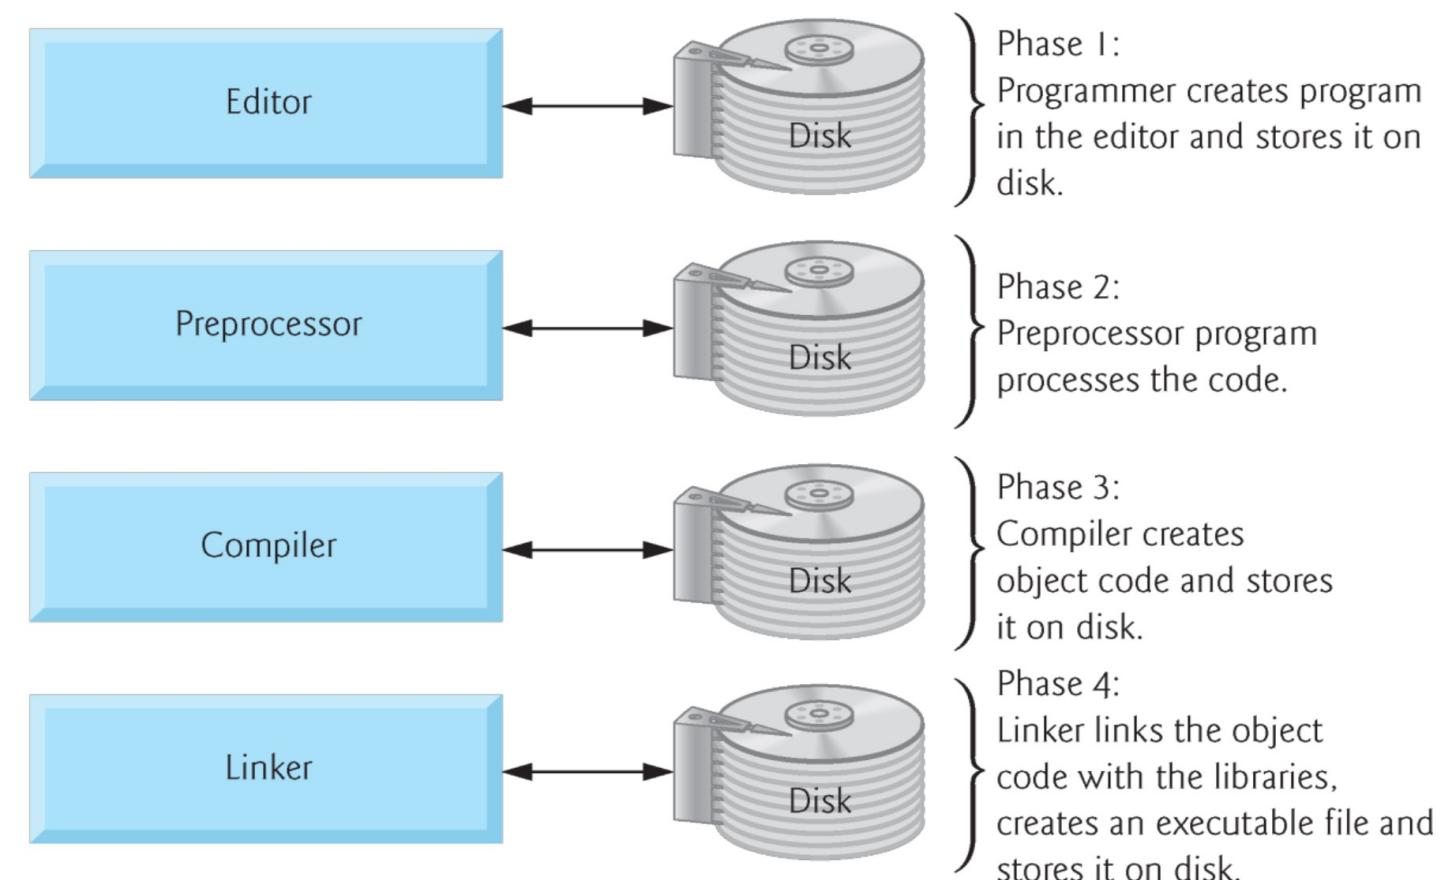
# Next Class

- C programming
- The **gcc** compiler
- Datatypes
- User input/output

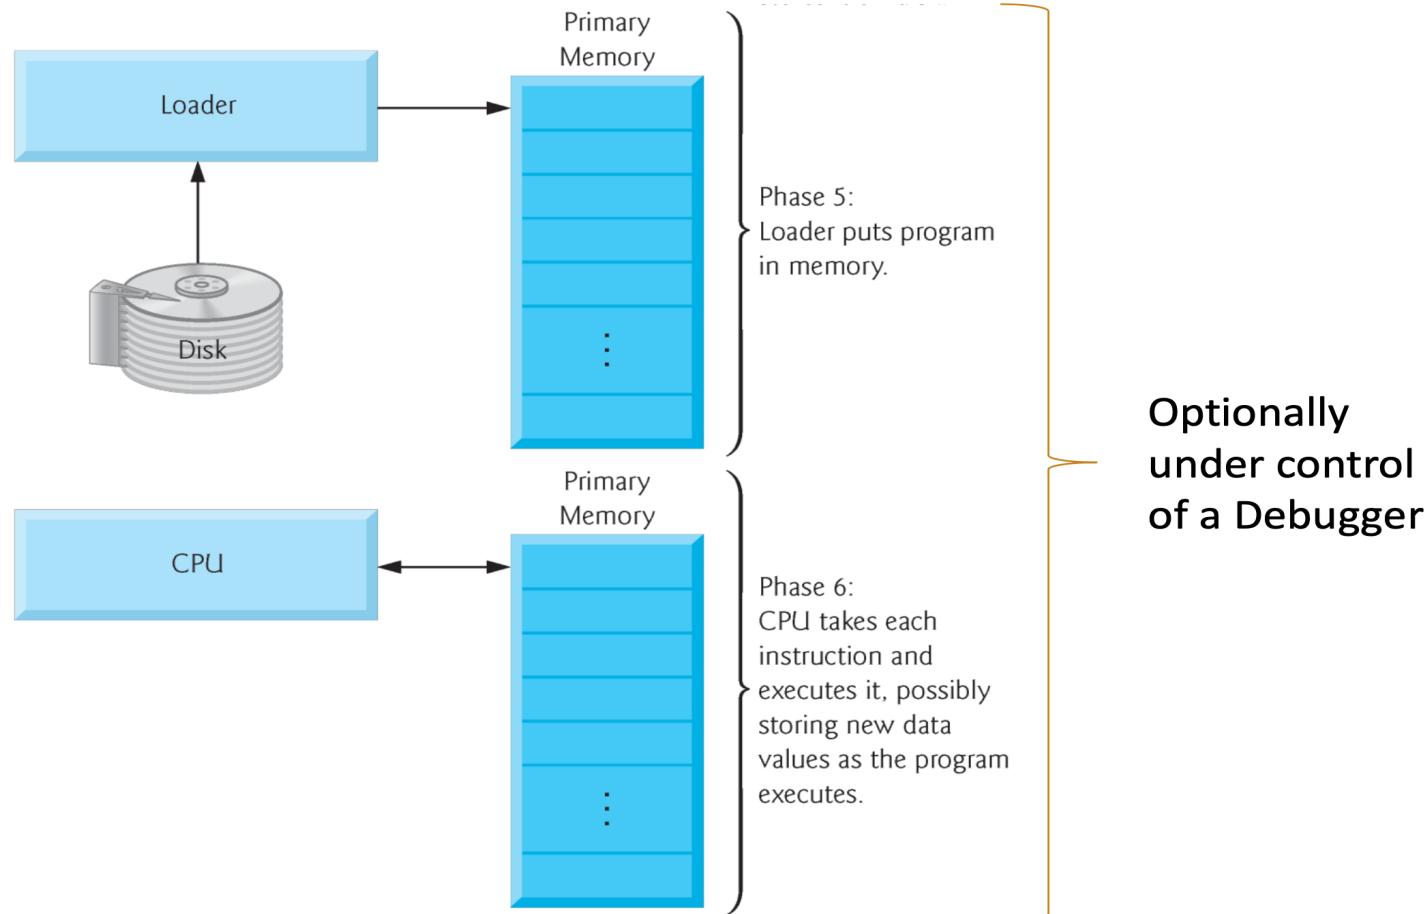
# History of C

- Developed at AT&T Bell Labs in early 1970s
- Unix also developed at Bell Labs
- All but core of Unix is in C
  - Linux and BSD are very similar to Unix
- Standardized by American National Standards Institute (ANSI)

# Phases of a C program



# Next Class



# Compiling a C program

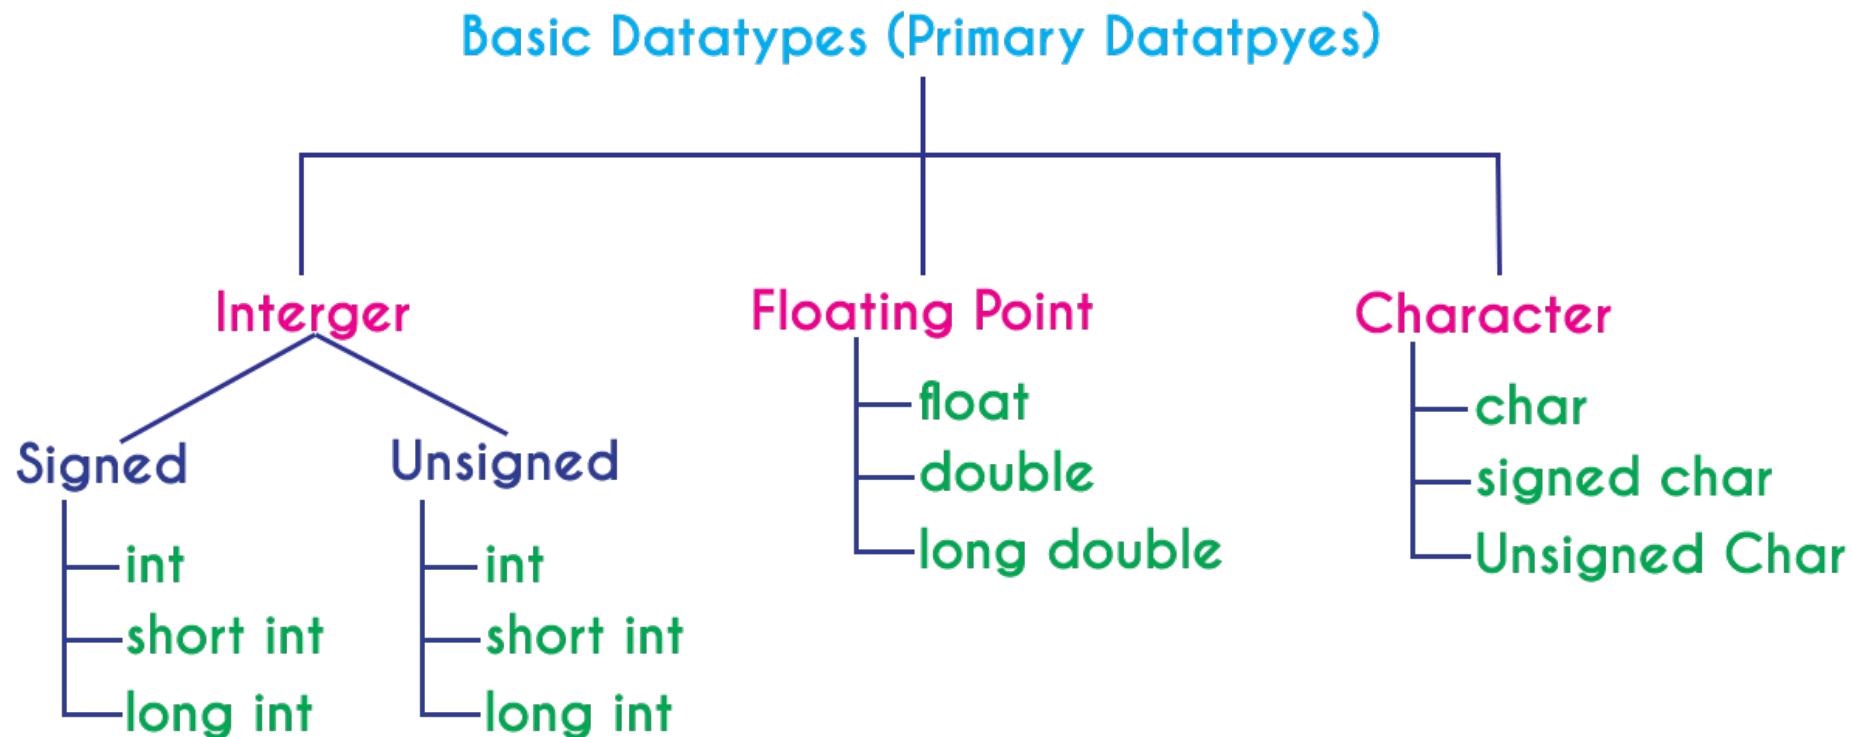
- To compile a C program we use **gcc**
  - `gcc hello_world.c`
  - `gcc -o hello_exec hello_world.c`
- Extra flags to see warnings:
  - **-Wall** – check for warnings
  - **-Wextra** – check for even more warnings
  - **-Werror** – treat warnings as errors
- Additional flags (not comprehensive):
  - **-S** (assembly), **-E** (preprocess only), **-c** (compile only), **-m32** (compile for 32 bit architecture)

# Hello World

```
#include <stdio.h>

int main()
{
    printf("Hello World!\n");
    return 0;
}
```

# Basic datatypes



# Basic datatypes

| Type           | Storage size                         | Value range  |
|----------------|--------------------------------------|--|
| char           | 1 byte                               | -128 to 127 or 0 to 255                              |
| unsigned char  | 1 byte                               | 0 to 255   |
| signed char    | 1 byte                               | -128 to 127  |
| int            | 2 or 4 bytes                         | -32,768 to 32,767 or -2,147,483,648 to 2,147,483,647 |
| unsigned int   | 2 or 4 bytes                         | 0 to 65,535 or 0 to 4,294,967,295                    |
| short          | 2 bytes                              | -32,768 to 32,767                                    |
| unsigned short | 2 bytes                              | 0 to 65,535  |
| long           | 8 bytes or<br>(4bytes for 32 bit OS) | -9223372036854775808 to 9223372036854775807          |
| unsigned long  | 8 bytes                              | 0 to 18446744073709551615                            |

# Basic datatypes

| Type        | Storage size | Value range            | Precision         |
|-------------|--------------|------------------------|-------------------|
| float       | 4 byte       | 1.2E-38 to 3.4E+38     | 6 decimal places  |
| double      | 8 byte       | 2.3E-308 to 1.7E+308   | 15 decimal places |
| long double | 10 byte      | 3.4E-4932 to 1.1E+4932 | 19 decimal places |

# Variables and datatypes examples

```
#include <stdio.h>

int integerVar = 42;
float floatVar = 3.14;

int main()
{
    double doubleVar = 2.71828;
    char charVar = 'A';

    printf("Integer: %d\n", integerVar);
    printf("Float: %f\n", floatVar);
    printf("Double: %lf\n", doubleVar);
    printf("Character: %c\n", charVar);

    return 0;
}
```

# Specifiers

| <b><i>specifier</i></b> | <b>Output</b>   | <b>Example</b> |
|-------------------------|---|----------------|
| d or i                  | Signed decimal integer  | 392            |
| u                       | Unsigned decimal integer  | 7235           |
| o                       | Unsigned octal  | 610            |
| x                       | Unsigned hexadecimal integer  | 7fa            |
| X                       | Unsigned hexadecimal integer (uppercase)  | 7FA            |
| f                       | Decimal floating point, lowercase   | 392.65         |
| F                       | Decimal floating point, uppercase   | 392.65         |
| e                       | Scientific notation (mantissa/exponent), lowercase  | 3.9265e+2      |
| E                       | Scientific notation (mantissa/exponent), uppercase  | 3.9265E+2      |
| g                       | Use the shortest representation: %e or %f   | 392.65         |
| G                       | Use the shortest representation: %E or %F   | 392.65         |
| a                       | Hexadecimal floating point, lowercase   | -0xc.90fep-2   |
| A                       | Hexadecimal floating point, uppercase   | -0XC.90FEP-2   |
| c                       | Character   | a              |
| s                       | String of characters  | sample         |
| p                       | Pointer address   | b8000000       |
| n                       | Nothing printed.<br>The corresponding argument must be a pointer to a signed int.<br>The number of characters written so far is stored in the pointed location. |                |
| %                       | A % followed by another % character will write a single % to the stream.  | %              |

# Constant variables

- C has two distinct ways of creating constants
  - `#define` preprocessor directive
  - `const` keyword
- Modern programs prefer `const`
  - Better code readability
  - Easier to debug
  - Follow the scope bounds

# Today's lecture

- Setting up SSH keys
- VS Code remote development

# SSH key generation

- AWS will generate a fresh key pair when you create a new instance
- You can add additional key pairs if you'd like to give other individuals access to your instance
- This generates a key pair, a public and a private key.
- The key must be in the `~/ .ssh/` directory
  - For windows it should be `C:\User\ .ssh\`
- The private key must have correct permissions:
  - `chmod 400 <path to private key>`

# VS Code remote development

- Press **Cmd (Ctrl for windows) + Shift + P**
- Search for Remote SSH add new agent
- Type in **ssh <username>@<host>** -- default username is ubuntu
- Open config file and add the following lines:

Host <IP of Instance>

HostName <IP of Instance>

User <username>

IdentityFile <Path to private key> (~/.ssh/key\_name.pem)

# Today's lecture

- Continuation of C discussion
  - Standard libraries
  - Operators
  - Loops and Conditionals
  - Functions
  - Libraries
  - Scope

# Standard libraries, and headers

- A header file contains function headers that allow your program to reference the library's functions
- The preprocessor combines these headers with your source program
- The linker hooks these references to the compiled object code of the library

| Library Header File | Functions  |
|---------------------|--|
| <b>stdio.h</b>      | Input and output, including file I/O                                     |
| <b>stdlib.h</b>     | Arithmetic, random numbers, terminating execution, sorting and searching |
| <b>math.h</b>       | Advanced math (trig)   |
| <b>string.h</b>     | String processing  |

# User input

- When we call **scanf** to read input from the user, we specify:
  - what type of data is being read,
  - a variable pointer into which we store the value.

```
#include <stdio.h>

int main()
{
    int num;

    printf("Enter an integer: ");
    scanf("%d", &num);
    printf("You entered: %x\n", num);

    return 0;
}
```

# Operators in C

| Operator symbol | Operation         | Example usage | Operator symbol | Operation   | Example usage |
|-----------------|-------------------|---------------|-----------------|-------------|---------------|
| *               | multiplication    | $x * y$       | $\sim$          | bitwise NOT | $\sim x$      |
| /               | division          | $x / y$       | $\&$            | bitwise AND | $x \& y$      |
| %               | integer remainder | $x \% y$      | $ $             | bitwise OR  | $x   y$       |
| +               | addition          | $x + y$       | $\wedge$        | bitwise XOR | $x \wedge y$  |
| -               | subtraction       | $x - y$       | $\ll$           | left shift  | $x \ll y$     |
|                 |                   |               | $\gg$           | right shift | $x \gg y$     |

| Operator symbol | Operation             | Example usage | Operator symbol | Operation   | Example usage |
|-----------------|-----------------------|---------------|-----------------|-------------|---------------|
| >               | greater than          | $x > y$       | !               | logical NOT | $!x$          |
| $\geq$          | greater than or equal | $x \geq y$    | $\&\&$          | logical AND | $x \&\& y$    |
| <               | less than             | $x < y$       | $  $            | logical OR  | $x    y$      |
| $\leq$          | less than or equal    | $x \leq y$    |                 |             |               |
| $\equiv$        | equal                 | $x \equiv y$  |                 |             |               |
| $\neq$          | not equal             | $x \neq y$    |                 |             |               |

# Conditionals and loops

|             |  |  |
|-------------|--|--|
| Conditional | <pre>if (a &lt; 1) {     // do something } else {     // do something else }</pre> | A condition determines which alternative to execute. |
| Iterative   | <pre>for (int i=1; i&lt;10; i++) {     // iterate }</pre>                          | Repeats statement as long as a condition is true.    |

# Class exercise

- Write a C program that calculate the area of a circle by prompting the user to input a radius value
- Your program should have the following requirements:
  - It must define pi as a constant with global scope.
  - Make sure to pick appropriate data types for your radius and area
  - Make use of the pow function to square the radius
  - If they user enters a negative radius value the program prints out an error and quits
- To install GCC on your VMs run the following command

```
sudo apt update  
sudo apt install gcc  
Check gcc --version (Should be 11.04)
```



bit.ly/297-1

# Today's lecture

- Functions
- Scope
- Libraries

# Functions

- Header: Everything before the first brace.
- Body: Everything between the braces.
- Type: Type of the value returned by the function.
- Parameter List: A list of identifiers that provide information for use within the body of the function.

```
type function_name ( parameter type list )
{
    declarations
    statements

    return type
}
```

# Examples

```
long long factorial(int n)
{
    if(n == 0 || n == 1)
    {
        return 1;
    }

    return n * factorial(n-1);
}
```

```
bool isPrime(int n)
{
    if (n <= 1) {
        return false;
    }

    for(int i=2; i*i <= n; i++)
    {
        if (n % i == 0) {return false;}
    }
    return true;
}
```

# Scope

- A block is code within enclosing braces { ... }
- A variable is declared within a block is local to that block, and is visible until the end of that block.
- Typically, a local variable will be declared at the beginning of a function block and used within that function.
- A variable declared outside any block is global. It is visible throughout the entire program.

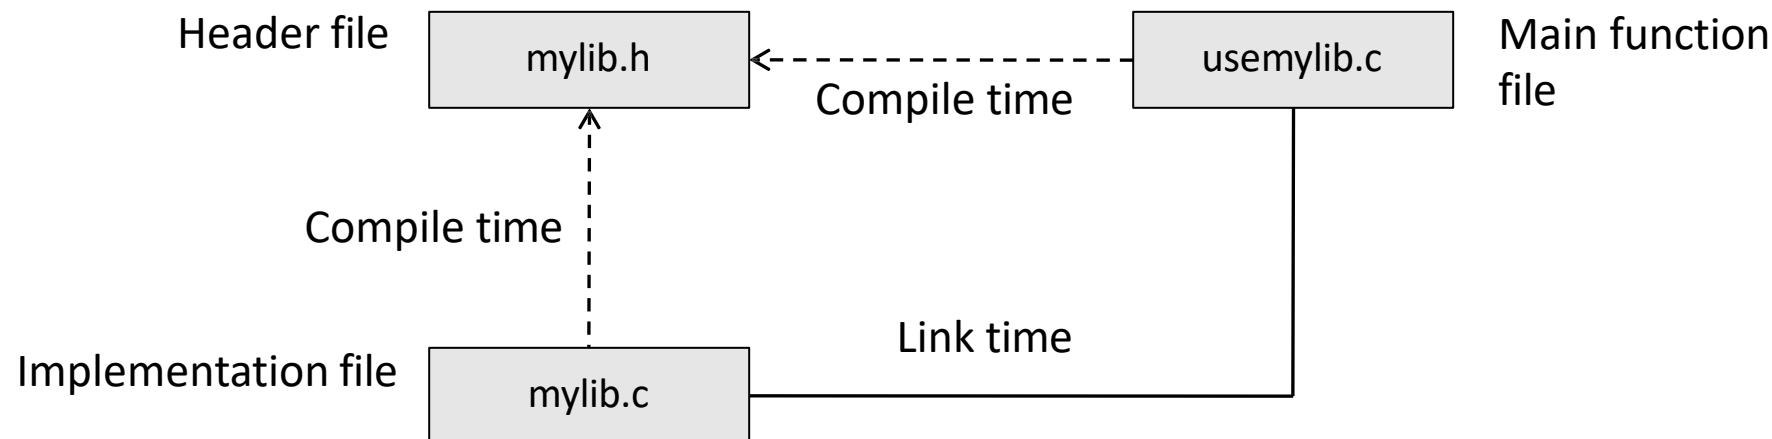
```
int count = 10;  
  
int main(void)  
{  
    int echo;  
    scanf("%d", &echo);  
  
    for (int i = 0; i < count; i++)  
    {  
        printf("%d\n", echo);  
    }  
}
```

# Placement of functions

- We can put them all in one file with the application; prototypes are at the beginning of the file, and we work top down
- But this requires recompiling the whole file after one small change
- Work can't be divided among many developers
- Functions can't be maintained independently of their clients

# Libraries

- Place the function prototypes for a module in a header file (.h extension)
- Place the completed function definitions for that module in an implementation file (.c)
- Compile the c. file, link it to a test driver, test it, and store it away



# Separate compilation and linking

- Save each implementation file with a .c extension
- Work bottom up, compile each .c file
  - `gcc -c <filename>.c`
- If no errors, link the object files
  - `gcc -o <main filename> <filename>.o ... <filename>.o`
- If no errors, run the program with `./<main filename>`

# An alternate approach – library archives

- An alternate approach is to create a library archive file (.a extension)
  - `ar rcs <output_file> <list of object files>`
- The convention is to prefix **lib** before the name of the library, and and the **.a** extension
- This can then be linked using the **-l** flag in gcc
  - `gcc <main_file> -L<path_to_lib> -l<name_of_lib> -o <output_file>`
  - When providing the name of the library, we do not need to provide the lib prefix or the .a extension

# Makefiles

- We use make files when there are many modules to keep track of
  - Which ones are up to date, which ones need to be recompiled?
- The UNIX **make** tool and makefiles help to manage this problem
  - When you run make, UNIX looks for a file named **makefile** or **Makefile** in the current working directory
- A makefile contains info about the dependencies among program modules and commands to compile and/or link them

# Makefile entries

**targets: dependencies  
commands**



A tab must go here

- Targets have specific dependencies
- Targets and dependencies are file names
- Commands are similar to those we've seen already e.g. (gcc ....)

# Checks for timestamps

- Makefile also checks for timestamps files are modified
- Compare time stamps of files
- If target file is older than any dependency file then execute commands

# Next class

- **Topic:** Pointers
- **Reading:** Chapter 5 of K&R

# Today's lecture

- Arrays
- Pointers

# An alternate approach – library archives

- An alternate approach is to create a library archive file (.a extension)
  - `ar rcs <output_file> <list of object files>`
- The convention is to prefix **lib** before the name of the library, and and the **.a** extension
- This can then be linked using the **-l** flag in gcc
  - `gcc <main_file> -L<path_to_lib> -l<name_of_lib> -o <output_file>`
  - When providing the name of the library, we do not need to provide the lib prefix or the .a extension

# Arrays

- A linear sequence of memory cells, each accessible by an integer index, ranging from 0 to the number of cells minus 1
- Like Python lists and Java array lists, C arrays support random (constant-time) access to the cells

# Declaring array variables

```
int integers[10]; // Cells for 10 ints
char letters[5]; // Cells for 5 characters
float matrix[3][2]; // A 3 by 2 matrix of floats

int max = 100; // The number of cells in the next array
double doubles[max]; // Cells for max doubles
```

**<element type> <variable> [<integer expression>]**

- Declares an array variable and reserves memory for a definite number of cells
  - use more [] for more dimensions
- Each of which can store an element of the given type (arrays are statically typed)
- Each cell initially contains garbage!

# Accessing array values

```
int i, max = 10, array[max] //Declare variables  
  
for (i = 0; i < max; i++)  
    array[i] = i + 1; //Store 1..max in the array  
  
for (i = 0; i < max; i++)  
    printf("%d\n", array[i]); //Print the contents
```

- Here we initialize all of the array's cells and visit them all
- The loop is pretty standard

# Arrays and functions

```
int minIndex(int array[], int length)
{
    int probe, minPos = 0;

    for (probe = 1; probe < length; probe++)
    {
        if (array[probe] < array[minPos])
            minPos = probe;
    }

    return minPos;
```

- An array can be passed as an argument to a function
- The function can access or replace values in the array cells

# Class exercise

```
int number, length = 0, max = 10, array[max];  
  
printf("Enter a number or 0 to stop: ");  
scanf("%d", &number);  
  
while (number)  
{  
    array[length] = number;  
    length++;  
  
    printf("Enter a number or 0 to stop: ");  
    scanf("%d", &number);  
}
```

Observe this code, and identify any errors/potential correctness issues.

[bit.ly/csci297-f23-2](https://bit.ly/csci297-f23-2)

# Parameter passing

- Parameters of basic types (char, int, float, and double) are passed by value (a copy of the value of the argument is placed in temporary storage on the runtime stack)
- A copy of an array's base address is also placed in temporary storage, so its cells can still be accessed or modified

# Pointers

- C allows the programmer to access the address of a memory location and use it as a data value
- Pointer = address of a variable (or any memory location)
  - Allows indirect access to variables – for example, allow a function to change a variable in the caller.
  - Pointer-based data structures that can grow or shrink to meet needs.
  - Pointers are also used to refer to dynamically allocated storage (from the system heap)

# Visualization

Name: ptrNumber (int\*)

Address: 0x????



**0x 22EE (&number)**

An int *pointer variable* contains  
a *memory address* pointing to  
an int value

Name: number (int)  
Address: 0x22EE (&number)

88

An int variable contains an int  
value

# Declaring pointers

- C allows us to declare a variable that holds a pointer value.
- We need to specify the type of data to which the pointer points to.
- Add `*` operator to the type name to declare a pointer to that type.
  - `int * P;`
  - `float *Q;`
  - `int *A[5];`
- P is a pointer of type int, Q is a pointer of type float, A is an array of 5 pointers of int
- **Common Confusion:** The use of whitespaces in declaration is irrelevant

# The address operator

- The & operator can be used before any variable to get its memory address

```
int main()
{
    int V = 5;

    printf ("The value of V: %d\n", V);
    printf ("The address of V: %p\n", &V);

}
```

# Assigning pointer values

```
int * P; //declare a pointer P of type int, int V = 5;  
  
printf ("The value of V: %d\n", V);  
  
printf ("The address of V: %p\n", &V);  
  
P = &V; //Pointer now points to the memory location of V  
  
printf ("The memory location P points to: %p\n", P);
```

# Dereferencing pointers

- The dereference operator (\*) is used to obtain the value of the variable a pointer points to.
- Done by adding \* before the pointer variable

```
int * P;  
int V = 5;  
P = &V;  
  
printf ("The value stored at P: %d\n", *P);
```

- Common Confusion: The dereference operator (\*) is different from the asterisk used to declare a pointer.

# Next class

- **Topic:** Dynamic memory / Heap
- **Reading:** Chapter 5 of K&R

# Today's lecture

- Assignment 1
- Pointers (Recap)
- Dynamic Memory

# Pointers

- C allows the programmer to access the address of a memory location and use it as a data value
- Pointer = address of a variable (or any memory location)
  - Allows indirect access to variables – for example, allow a function to change a variable in the caller.
  - Pointer-based data structures that can grow or shrink to meet needs.
  - Pointers are also used to refer to dynamically allocated storage (from the system heap)

# Visualization

Name: ptrNumber (int\*)

Address: 0x????



**0x 22EE (&number)**

An int *pointer variable* contains  
a *memory address* pointing to  
an int value

Name: number (int)  
Address: 0x22EE (&number)



88

An int variable contains an int  
value

# Declaring pointers

- C allows us to declare a variable that holds a pointer value.
- We need to specify the type of data to which the pointer points to.
- Add `*` operator to the type name to declare a pointer to that type.
  - `int * P;`
  - `float *Q;`
  - `int *A[5];`
- P is a pointer of type int, Q is a pointer of type float, A is an array of 5 pointers of int
- **Common Confusion:** The use of whitespaces in declaration is irrelevant

# The address operator

- The & operator can be used before any variable to get its memory address

```
int main()
{
    int V = 5;

    printf ("The value of V: %d\n", V);
    printf ("The address of V: %p\n", &V);

}
```

# Assigning pointer values

```
int * P; //declare a pointer P of type int, int V = 5;  
  
printf ("The value of V: %d\n", V);  
  
printf ("The address of V: %p\n", &V);  
  
P = &V; //Pointer now points to the memory location of V  
  
printf ("The memory location P points to: %p\n", P);
```

# Dereferencing pointers

- The dereference operator (\*) is used to obtain the value of the variable a pointer points to.
- Done by adding \* before the pointer variable

```
int * P;  
int V = 5;  
P = &V;  
  
printf ("The value stored at P: %d\n", *P);
```

- Common Confusion: The dereference operator (\*) is different from the asterisk used to declare a pointer.

# NULL Pointer

- Sometimes we want to create a pointer that points to nothing.
- This is known as a null pointer
  - `int *ptr = NULL;`
- Why would we want to do this?
- Initialize unused pointers to NULL. If you don't initialize, the pointer will point to a random memory location, which can lead to strange and interesting bugs.

# Pointers and functions

```
#include <stdio.h>

void swap(int *a, int *b)
{
    int temp = *a;
    *a = *b;
    *b = temp;
}

int main()
{
    int x = 10, y = 20;

    printf("Before swap: x = %d, y = %d\n", x, y);
    swap(&x, &y);
    printf("After swap: x = %d, y = %d\n", x, y);

    return 0;
}
```

# Dynamic memory

- C allows the program to allocate memory in real time on the heap
- These memory locations are referenced by pointer
- Benefits
  - No need to decide beforehand the size of data structures e.g (arrays)
  - Allows working with variable input data
  - Also efficient and saves memory as we only allocate only what is needed.
- C functions for dynamic memory allocation:
  - `malloc` allocates a block of storage from the heap
  - `free` returns storage to the heap

# Malloc

- Allocates requested number of bytes
- Returns a void pointer to the first byte of the allocated space
- Malloc allocates a contiguous chunk of memory
  - `ptr= (int*) malloc(100 *sizeof(int));`
- int is of 4 bytes, so the total bytes allocated are 400

# Free

- Use the free command to free the allocated memory
  - `free(ptr);`
- With C we have to manually clear memory we allocate
- Question: What if we don't free the memory?

# Next class

- **Topic:** File I/O and Structs
- **Reading:** Chapter 6 and 7 of K&R

# Today's lecture

- Dynamic Memory

# Dynamic memory

- C allows the program to allocate memory in real time on the heap
- These memory locations are referenced by pointer
- Benefits
  - No need to decide beforehand the size of data structures e.g (arrays)
  - Allows working with variable input data
  - Also efficient and saves memory as we only allocate only what is needed.
- C functions for dynamic memory allocation:
  - `malloc` allocates a block of storage from the heap
  - `free` returns storage to the heap

# Malloc

- Allocates requested number of bytes
- Returns a void pointer to the first byte of the allocated space
- Malloc allocates a contiguous chunk of memory
  - `ptr= (int*) malloc(100 *sizeof(int));`
- int is of 4 bytes, so the total bytes allocated are 400

# Free

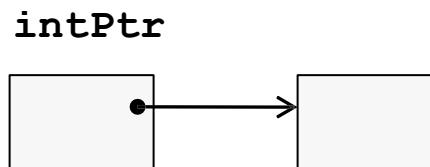
- Use the free command to free the allocated memory
  - `free(ptr);`
- With C we have to manually clear memory we allocate
- Question: What if we don't free the memory?

# Declare and Initialize with malloc

- **malloc** expects an integer representing the size of the storage needed
- The **sizeof** function takes a type as an argument and returns the number of bytes needed for a value of that type
- **malloc** returns a pointer to the first byte of the storage

```
#include <stdlib.h>

int *intPtr = malloc(sizeof(int));
```



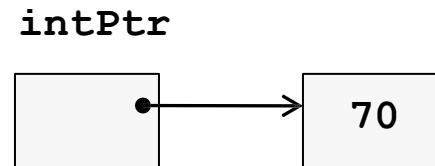
The variable now points to a new storage area for an **int**, but this storage contains garbage

# Access via Dereference with \*

```
#include <stdlib.h>

int *intPtr = malloc(sizeof(int));
(*intPtr) = 70;
printf("%d\n", *intPtr);
```

- The \* operator accesses the cell pointed to
- The pointer must point to a cell
- \* has a lower precedence than =



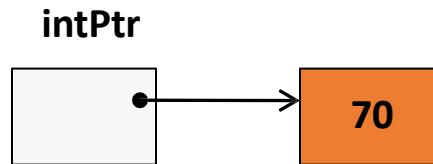
The variable now points to storage that has a program-supplied value

# Return the Storage with free

```
#include <stdlib.h>

int *intPtr = malloc(sizeof(int));
(*intPtr) = 70;
printf("%d\n", *intPtr);
free(intPtr);
```

- The storage is still accessible, but might be taken for other applications (also called a dangling pointer reference)



The variable still points to storage, which might be used by others

# Reset the Pointer to NULL

```
#include <stdlib.h>

int *intPtr = malloc(sizeof(int));
(*intPtr) = 70;
printf("%d\n", *intPtr);
free(intPtr);
intPtr = NULL;
```

- Clean up by setting the pointer to **NULL** or to other storage

0

70

The variable now has no dangling pointer

# Dynamic memory as arrays

```
int *arr = (int*)malloc(10 * sizeof(int));  
  
// Populate the array with numbers 1 through 10.  
for (int i = 0; i < 10; i++)  
{  
    arr[i] = i + 1;  
}  
  
int *temp = (int*)realloc(arr, 20 * sizeof(int));  
  
// Check if the memory reallocation was successful.  
if (temp == NULL) {return 1;}  
  
arr = temp; // Update the pointer to the new memory location.  
  
// Don't forget to free the memory when done.  
free(arr);  
  
return 0;
```

# Structs

- An issue with arrays is you can only store one data type.
- C supports data structures called structs that can support a combination of data types such as ints, floats and char types.
- A struct is a data type that includes one or more named data fields
- Good for containing data of different types
- Like classes in Python or Java, but without the methods

# Typedef

- C allows you to give data types custom names or aliases
- This is done using the `typedef` keyword. Syntax:
- `MYINT` is the new alias. It is convention to use uppercase letters, but lowercase is allowed
- `typedef` can also be used with custom data types (eg. Structs)

```
typedef int MYINT;  
  
MYINT a; // a is of an integer type MYINT  
  
b[10] // array of 10 ints
```

# Next class

- **Topic:** Valgrind and GDB – Bring Laptops

# Today's lecture

- Structs wrap-up
- Valgrind

# Struct Example

```
// Define a struct to hold a dynamic array and its size
typedef struct {

    int *array; // Pointer to the first element of the dynamic array
    int size, count; // Number of elements in the array
} DynamicArray;
```

# Initializing the array struct

```
// Function to initialize the dynamic array with a given size
DynamicArray initializeArray(int size) {

    DynamicArray arr;

    arr.array = (int *)malloc(size * sizeof(int));

    if (arr.array == NULL)
    {
        fprintf(stderr, "Memory allocation failed!\n");
        exit(EXIT_FAILURE);
    }
    arr.size = size;
    arr.count = 0; // Initialize the count to zero

    return arr;
}
```

# Adding an element

```
// Function to add an element to the dynamic array
int addElement(DynamicArray *arr, int element) {

    if (arr->count < arr->size) {

        arr->array[arr->count] = element;
        arr->count++;

        return 0; // Success
    }
    return -1; // Failure: array is full
}
```

# Freeing the array

```
// Function to free the dynamic array
void freeArray(DynamicArray *arr) {

    free(arr->array);

    arr->array = NULL;

    arr->size = 0;

    arr->count = 0;
}
```

# Main function argument vector

- The argument vector is an array of strings that are gathered up when a C program is run at the command prompt
- The first string is the name of the command
- The remaining strings are the arguments, if any
- The argument vector and its length are optional formal parameters of the main function

# Example

```
/*
Prints contents of the argument vector (the name of the program
and any command-line arguments).
*/
#include <stdio.h>

int main(int length, char *argVec[])
{
    int i;
    for (i = 0; i < length; i++)
        printf("Argument %d: %s\n", i, argVec[i]);
}
```

# Valgrind

- Memory management, leak detection, and analysis tool.
- Detect memory leaks and errors in C/C++ programs.
- Valgrind flags errors that don't appear without valgrind
- Installing Valgrind:  
`sudo apt-get install valgrind  
valgrind --version`
- Running Valgrind
  - `valgrind ./program`

# Valgrind

- When to use?
  - Dealing with memory (especially dynamic memory allocation)
  - Whenever bugs occur. Get instant feedback about what the bug is, where it occurred, and why.
- When not to use?
  - Program contains no invalid reads and writes and no leaked memory
  - If the test case is inherently slow, then this is not a good choice

# Example run

```
#include <stdlib.h>

int main() {
    int *ptr = malloc(sizeof(int));
return 0;
}
```

```
ubuntu@ip-172-31-29-123:~/vg$ valgrind ./a.out
==2081== Memcheck, a memory error detector
==2081== Copyright (C) 2002-2017, and GNU GPL'd, by Julian Seward et al.
==2081== Using Valgrind-3.18.1 and LibVEX; rerun with -h for copyright info
==2081== Command: ./a.out
==2081==
==2081==
==2081== HEAP SUMMARY:
==2081==     in use at exit: 4 bytes in 1 blocks
==2081==   total heap usage: 1 allocs, 0 frees, 4 bytes allocated
==2081==
==2081== LEAK SUMMARY:
==2081==     definitely lost: 4 bytes in 1 blocks
==2081==     indirectly lost: 0 bytes in 0 blocks
==2081==     possibly lost: 0 bytes in 0 blocks
==2081==     still reachable: 0 bytes in 0 blocks
==2081==           suppressed: 0 bytes in 0 blocks
==2081== Rerun with --leak-check=full to see details of leaked memory
==2081==
==2081== For lists of detected and suppressed errors, rerun with: -s
==2081== ERROR SUMMARY: 0 errors from 0 contexts (suppressed: 0 from 0)
```

# Recommended flags

- Recommended flags

```
valgrind --leak-resolution=high --leak-check=full --show-reachable=yes --  
track-fds=yes -track-origins=yes ./myProgram
```

- Feel free to lookup **man valgrind**

# Types of memory leaks

- Still Reachable
  - Block is still pointed at, programmer could go back and free it before exiting
- Definitely Lost
  - No pointer to the block can be found
- Indirectly Lost
  - Block is “lost” because the blocks that point to it are themselves lost
- Possibly Lost
  - Pointer exists but it points to an internal part of the memory block

# Class exercise

- View the code provided in slack, run is using valgrind and post your observations at this link:

<https://bit.ly/csci297-f23-3>

# Today's lecture

- File I/O
- Introduction to linking

# Interacting with files

```
#include <stdio.h>

int main(){
    FILE *infile = fopen("test_file.c", "r");

    FILE *outfile = fopen("test.bin", "w");

    if (infile == NULL)
        {printf("The input file does not exist.\n");}
    else
    {
        processFile(infile, outfile);
        fclose(infile);
        fclose(outfile);
    }
}
```

# File I/O functions

| Function  | What it does  |
|---|---|
| <code>char* fgets(char* array,<br/>           int arraySize,<br/>           FILE *infile)</code>      | Returns <code>NULL</code> if end of file ( <code>EOF</code> ). Otherwise, loads the next line, including the newline and <code>null</code> characters, into <code>array</code> .<br><b>Note:</b> there must be storage for <code>array</code> ! |
| <code>int fputs(char* array,<br/>          FILE *outfile)</code>                                      | Returns <code>EOF</code> if the string cannot be written to <code>outfile</code> . Otherwise, writes the string to <code>outfile</code> . A newline is not automatically written.   |
| <code>fprintf(FILE *outfile,<br/>         char* formatString,<br/>         arg1, . . . , argn)</code> | Works just like <code>printf</code> with the standard output file (the terminal)  |

# stdin and stdout

```
#include <stdio.h>
#include <stdlib.h>

int main() {
    char filename[256]; // Buffer to hold the filename
    // Prompt the user to enter the filename

    printf("Enter the filename: ");
    if (fgets(filename, sizeof(filename), stdin) == NULL)
    {
        perror("Error reading filename");
        return 1;
    }

    // Open the file in read mode
    FILE *file = fopen(filename, "r");

    char buffer[256]; // Buffer to hold lines read from the file
    // Read lines from the file and write them to stdout

    while (fgets(buffer, sizeof(buffer), file) != NULL)
    {
        fputs(buffer, stdout);
    }

    // Close the file after reading
    fclose(file);

    return 0;
}
```

# Example C Program

**main.c**

```
int sum(int *a, int n);

int array[2] = {1, 2};

int main(int argc, char** argv)
{
    int val = sum(array, 2);

    return val;
}
```

**sum.c**

```
int sum(int *a, int n)
{
    int i, s = 0;

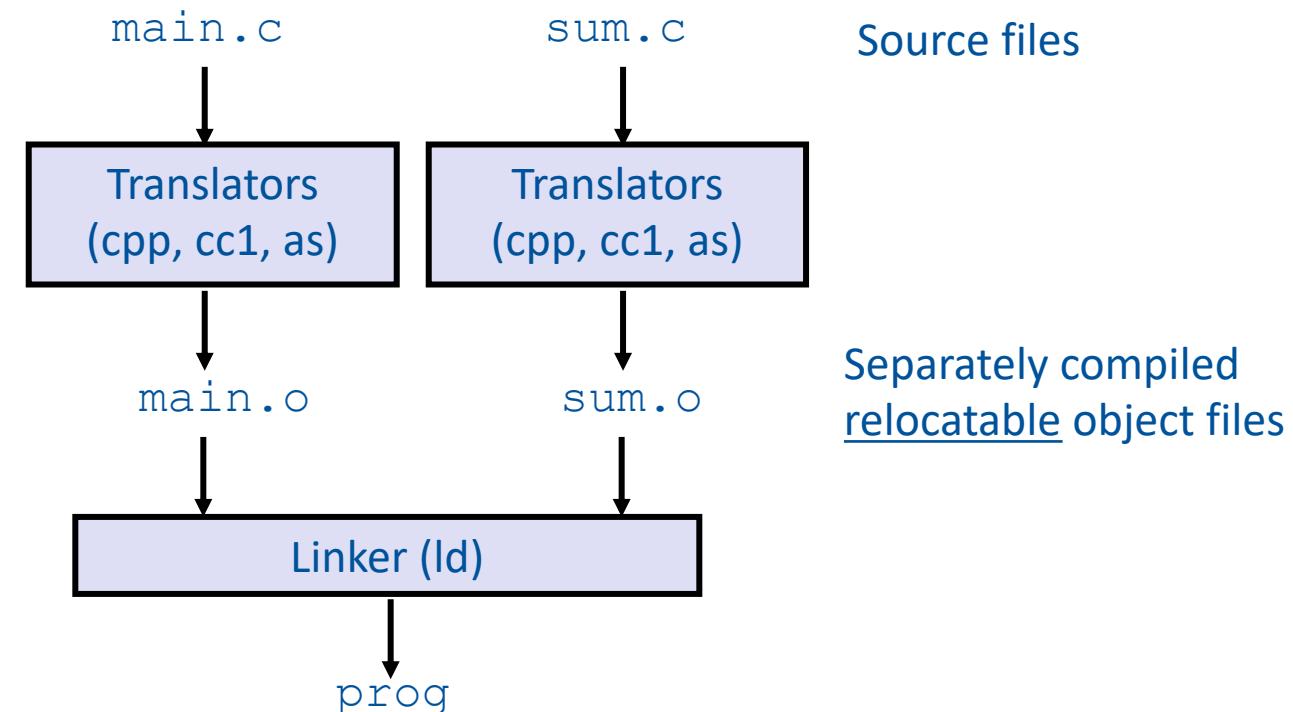
    for (i = 0; i < n; i++)
    {
        s += a[i];
    }

    return s;
}
```

# Linking

- Programs are translated and linked using a *compiler driver*:
  - `linux> gcc -Og -o prog main.c sum.c`
  - `linux> ./prog`

Fully linked executable object file  
(contains code and data for all functions  
defined in main.c and sum.c)



# Next Class:

- The ELF format and relocation
- Readings: 7.1 -7.7

# Today's lecture

- Linking steps
- The ELF format

# Why use linkers?

- Reason 1: Modularity
  - Program can be written as a collection of smaller source files, rather than one monolithic mass.
  - Can build libraries of common functions
- Reason 2: Time Efficiency
  - Time: Separate compilation
  - Change one source file, compile, and then relink
  - No need to recompile other source files, and can compile multiple files concurrently

# Why use linkers?

- Reason 3: Space Efficiency
  - Common functions can be aggregated into a single file
  - Option 1: Static Linking
    - Executable files and running memory images contain only the library code they actually use
  - Option 2: Dynamic Linking
    - Executable files contain no library code
    - During execution, single copy of library code can be shared across all executing processes

# Linker steps

- Step 1: Symbol resolution
  - Programs define and reference symbols (globals, functions, and some locals)
    - `void swap() { ... } /* define symbol swap */`
    - `swap(); /* reference symbol swap */`
    - `int *xp = &x; /* define symbol xp, reference x */`
  - Symbol definitions are stored in object file (by assembler) in *symbol table*.
    - Symbol table is an array of entries
    - Each entry includes name, size, and location of symbol.
  - **During symbol resolution step, the linker associates each symbol reference with exactly one symbol definition.**

# Identify the symbols

**main.c**

```
int sum(int *a, int n);

int array[2] = {1, 2};

int main(int argc, char** argv)
{
    int val = sum(array, 2);

    return val;
}
```

**sum.c**

```
int sum(int *a, int n)
{
    int i, s = 0;

    for (i = 0; i < n; i++)
    {
        s += a[i];
    }

    return s;
}
```

# Linker steps

- Step 2: Relocation
  - Merges separate code and data sections into single sections
  - Relocates symbols from their relative locations in the .o files to their final absolute memory locations in the executable.
  - Updates all references to these symbols to reflect their new positions.

# Three kinds of object files

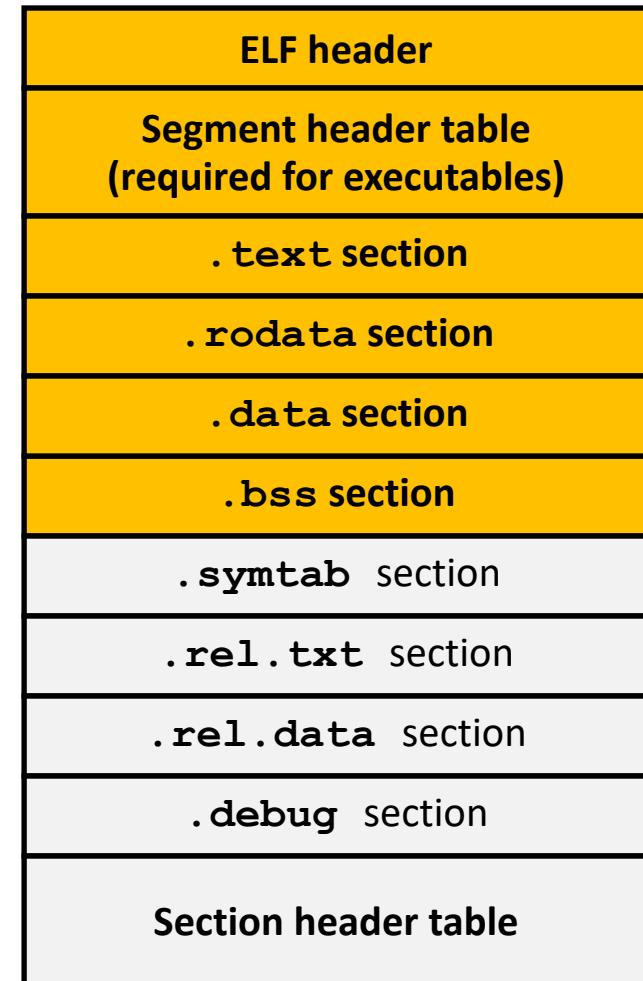
- Relocatable object file (.o file)
  - Contains code and data in a form that can be combined with other relocatable object files to form executable file.
  - Each .o file is produced from exactly one source (.c) file
- Executable object file (a.out file)
  - Contains code and data in a form that can be copied directly into memory and then executed.
- Shared object file (.so file)
  - Special type of relocatable object file that can be loaded into memory and linked dynamically, at either load time or run-time.
  - Called Dynamic Link Libraries (DLLs) by Windows

# Executable and linkable format

- Standard binary format for object files
- One unified format for
  - Relocatable object files (.o),
  - Executable object files (a.out)
  - Shared object files (.so)
- Generic name: ELF binaries

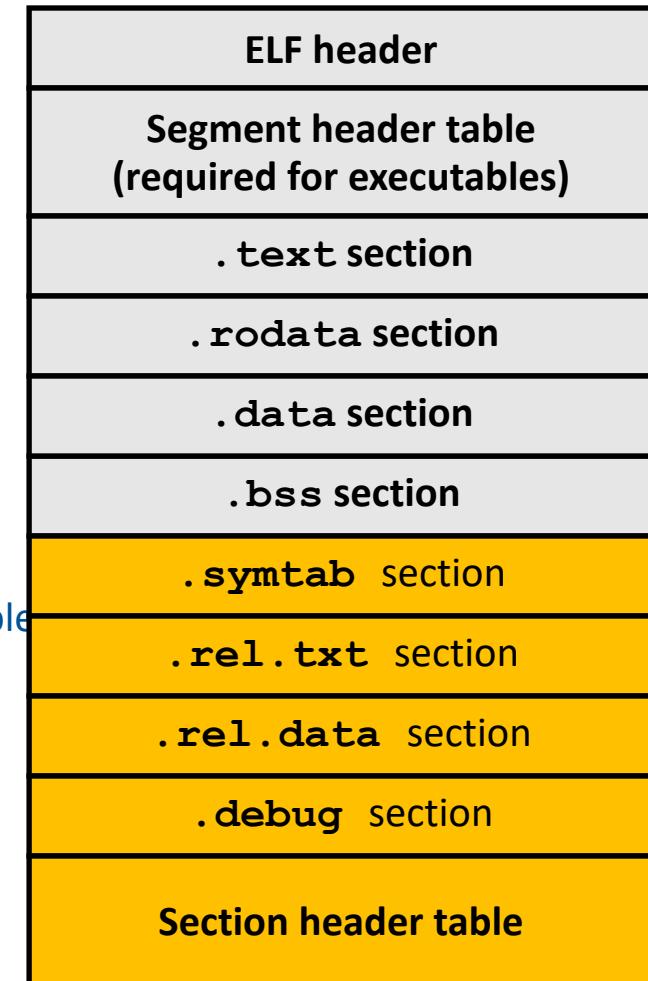
# ELF object file format

- **Elf header**
  - Word size, byte ordering, file type (.o, exec, .so), machine type, etc.
- **Segment header table**
  - Page size, virtual address memory segments (sections), segment sizes.
- **.text section**
  - Code
- **.rodata section**
  - Read only data: jump tables, string constants, ...
- **.data section**
  - Initialized global variables
- **.bss section**
  - Uninitialized global variables
  - “Block Started by Symbol”
  - “Better Save Space”
  - Has section header but occupies no space



# ELF object file format

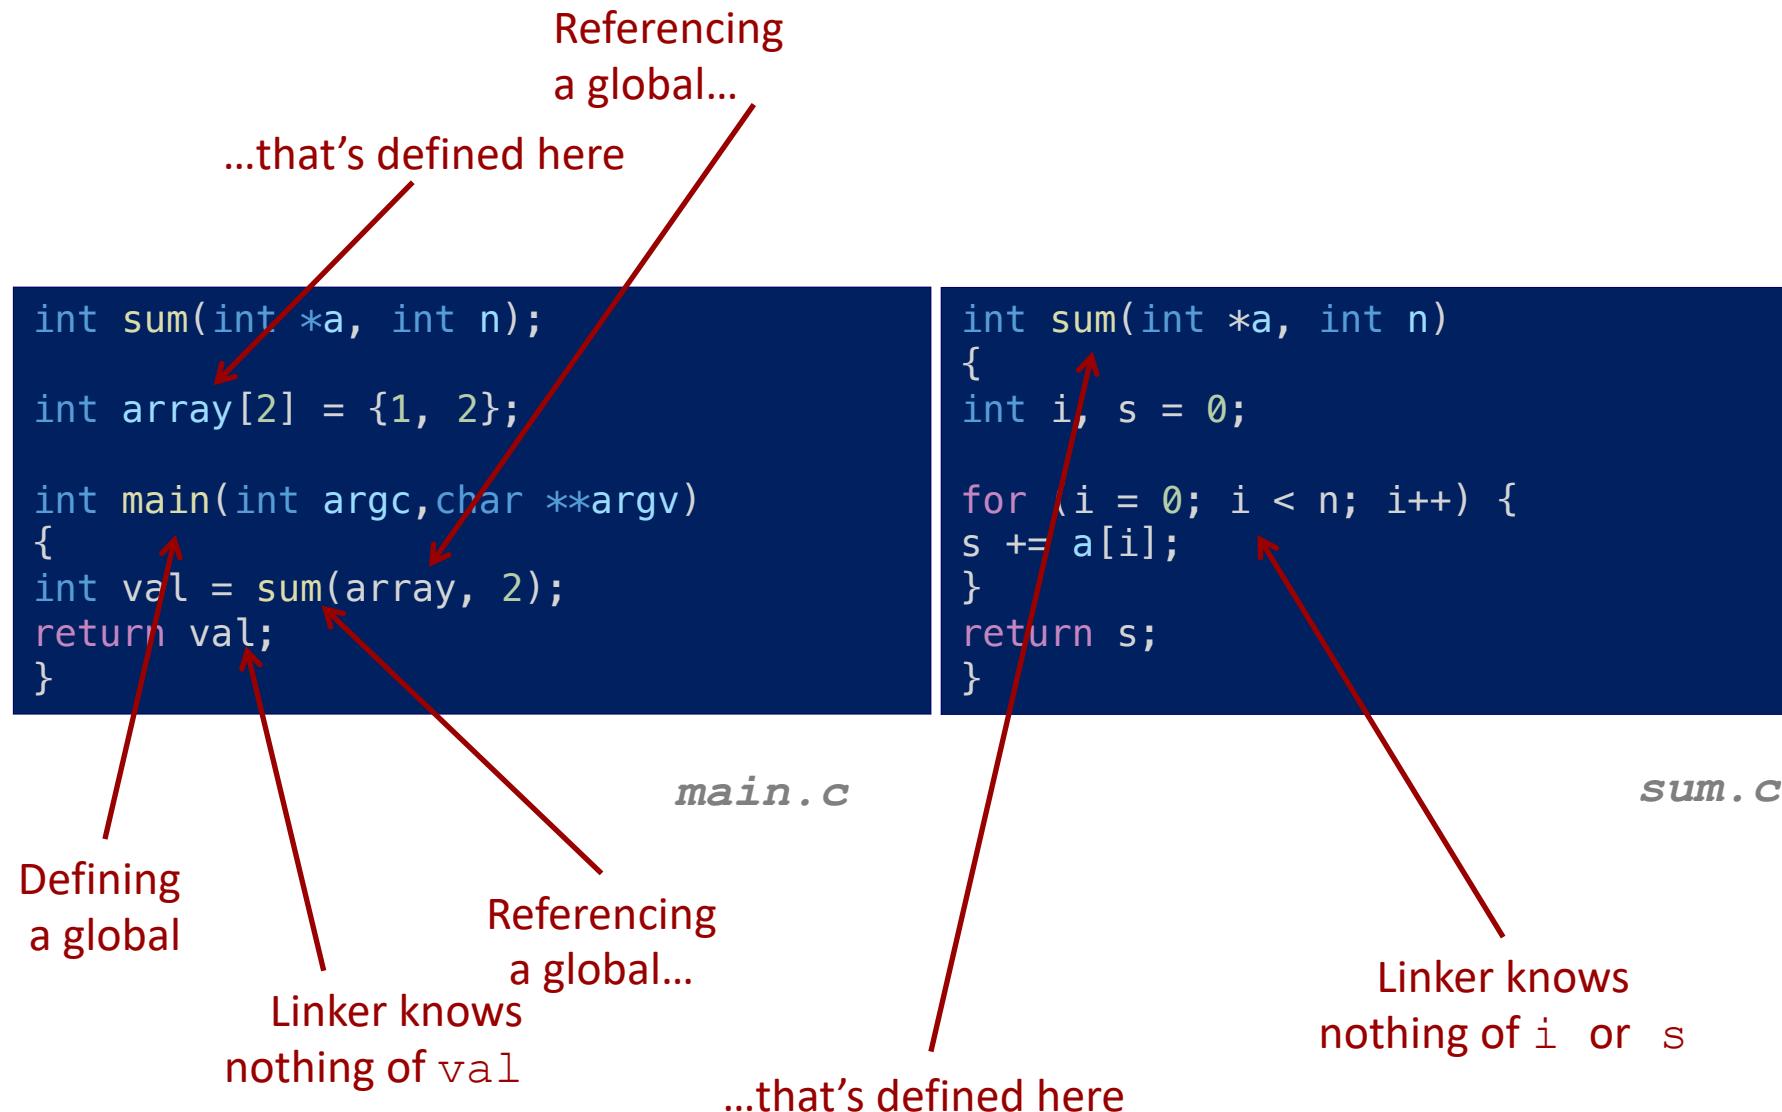
- **.syntab** section
  - Symbol table
  - Procedure and static variable names
  - Section names and locations
- **.rel.text** section
  - Relocation info for **.text** section
  - Addresses of instructions that will need to be modified in the executable
  - Instructions for modifying
- **.rel.data** section
  - Relocation info for **.data** section
  - Addresses of pointer data that will need to be modified in the merged executable
- **.debug** section
  - Info for symbolic debugging (`gcc -g`)
- **Section header table**
  - Offsets and sizes of each section



# Linker symbols

- Global symbols
  - Symbols defined by module m that can be referenced by other modules.
  - e.g., non-static C functions and non-static global variables.
- External symbols
  - Global symbols that are referenced by module m but defined by some other module.
- Local symbols
  - Symbols that are defined and referenced exclusively by module m.
  - e.g., C functions and global variables defined with the static attribute, and local static variables
  - **Local linker symbols are not local program variables**

# Symbol identification



# Symbol identification

- Which of the following names will be in the symbol table of symbols.o?

```
int incr = 1;
static int foo(int a)
{
    int b = a + incr;
    return b;
}

int main(int argc, char* argv[])
{
    printf("%d\n", foo(5));
    return 0;
}
```

- **incr**
- **foo**
- **a**
- **b**
- **argc**
- **argv**
- **main**
- **printf**
- **"%d\n"**

# Local symbols

- Local non-static C variables vs. local static C variables
  - Local non-static C variables: stored on the stack
  - Local static C variables: stored in either **.bss** or **.data**

```
static int x = 15;

int f() {
    static int x = 17;
    return x++;
}

int g() {
    static int x = 19;
    return x += 14;
}
```

Compiler allocates space in **.data** for each definition of **x**

Creates local symbols in the symbol table with unique names,  
e.g., **x**, **x.1721** and **x.1724**.

# Class exercise

```
int time;  
  
int foo(int x, int y);  
  
int bar(int a) {  
    int b = a + 1;  
    return b;  
}  
  
int main() {  
    printf("%d", bar(5));  
}
```

How many entries do we expect to see in the symbol table?

- A. 1
- B. 2
- C. 3
- D. 4
- E. > 4

<https://bit.ly/csci-297-f23-4>

# Next Class:

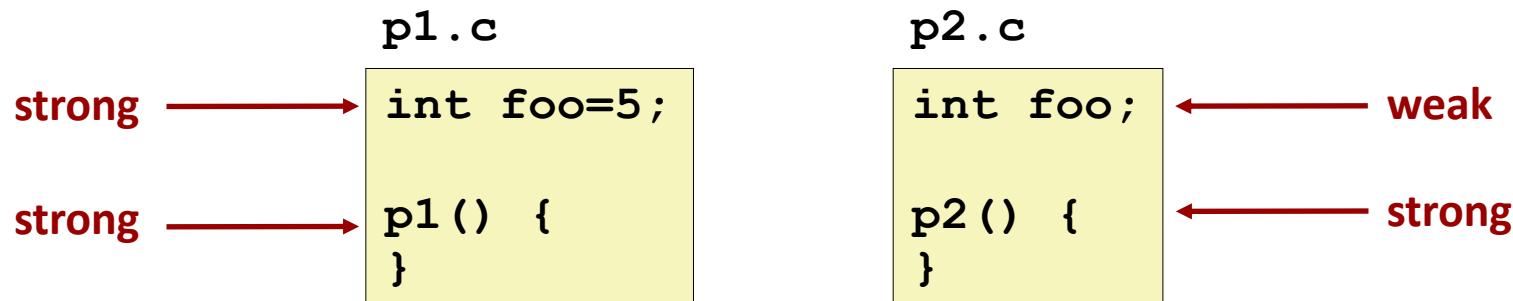
- Symbols resolution and relocation
- Readings: 7.1 -7.7

# Today's lecture

- Symbol resolution
- Linking with static libraries

# Duplicate symbols

- Program symbols are either strong or weak
  - Strong: procedures and initialized globals
  - Weak: uninitialized globals, or ones declared with specifier `extern`



# Symbol rules

- Rule 1: Multiple strong symbols are not allowed
  - Each item can be defined only once
  - Otherwise: Linker error
- Rule 2: Given a strong symbol and multiple weak symbols, choose the strong symbol
  - References to the weak symbol resolve to the strong symbol
- Rule 3: If there are multiple weak symbols, pick an arbitrary one
- DEMO!

# Practice questions

```
int x;  
p1() {}
```

```
p1() {}
```

**Link time error: two strong symbols (p1)**

```
int x;  
p1() {}
```

```
int x;  
p2() {}
```

**References to x will refer to the same uninitialized int. Is this what you really want?**

```
int x;  
int y;  
p1() {}
```

```
double x;  
p2() {}
```

**Writes to x in p2 might overwrite y!  
Evil!**

```
int x=7;  
int y=5;  
p1() {}
```

```
double x;  
p2() {}
```

**Writes to x in p2 might overwrite y!  
Nasty!**

```
int x=7;  
p1() {}
```

```
int x;  
p2() {}
```

**References to x will refer to the same initialized variable.**

# Global variables

- Avoid if you can
- Otherwise
  - Use **static** if you can
  - Initialize if you define a global variable
  - Use **extern** if you reference an external global variable
    - Treated as weak symbol
    - But also causes linker error if not defined in some file

# Use of extern

c1.c

```
#include "global.h"

int f() {
    return g+1;
}
```

global.h

```
extern int g;
int f();
```

c2.c

```
#include <stdio.h>
#include "global.h"

int g = 0;

int main(int argc, char argv[]) {
    int t = f();
    printf("Calling f yields %d\n", t);
    return 0;
}
```

# Packaging functions

- How to package functions commonly used by programmers?
  - Math, I/O, memory management, string manipulation, etc.
- Awkward, given the linker framework so far:
  - **Option 1:** Put all functions into a single source file
    - Programmers link big object file into their programs
    - Space and time inefficient
  - **Option 2:** Put each function in a separate source file
    - Programmers explicitly link appropriate binaries into their programs
    - More efficient, but burdensome on the programmer

# Static libraries

- Static libraries (**.a** archive files)
  - Concatenate related relocatable object files into a single file with an index (called an *archive*).
  - Enhance linker so that it tries to resolve unresolved external references by looking for the symbols in one or more archives.
  - If an archive member file resolves reference, link it into the executable.

# Next Class

- Relocation, Dynamic libraries
- **Readings:** 7.7 – 7.14

# Today's lecture

- Linking with static libraries
- Relocation
- Executables

# Packaging functions

- How to package functions commonly used by programmers?
  - Math, I/O, memory management, string manipulation, etc.
- Awkward, given the linker framework so far:
  - **Option 1:** Put all functions into a single source file
    - Programmers link big object file into their programs
    - Space and time inefficient
  - **Option 2:** Put each function in a separate source file
    - Programmers explicitly link appropriate binaries into their programs
    - More efficient, but burdensome on the programmer

# Static libraries

- Static libraries (**.a** archive files)
  - Concatenate related relocatable object files into a single file with an index (called an *archive*).
  - Enhance linker so that it tries to resolve unresolved external references by looking for the symbols in one or more archives.
  - If an archive member file resolves reference, link it into the executable.

# Common C libraries

## **libc.a** (the C standard library)

- 4.6 MB archive of 1496 object files.
- I/O, memory allocation, signal handling, string handling, data and time, random numbers, integer math

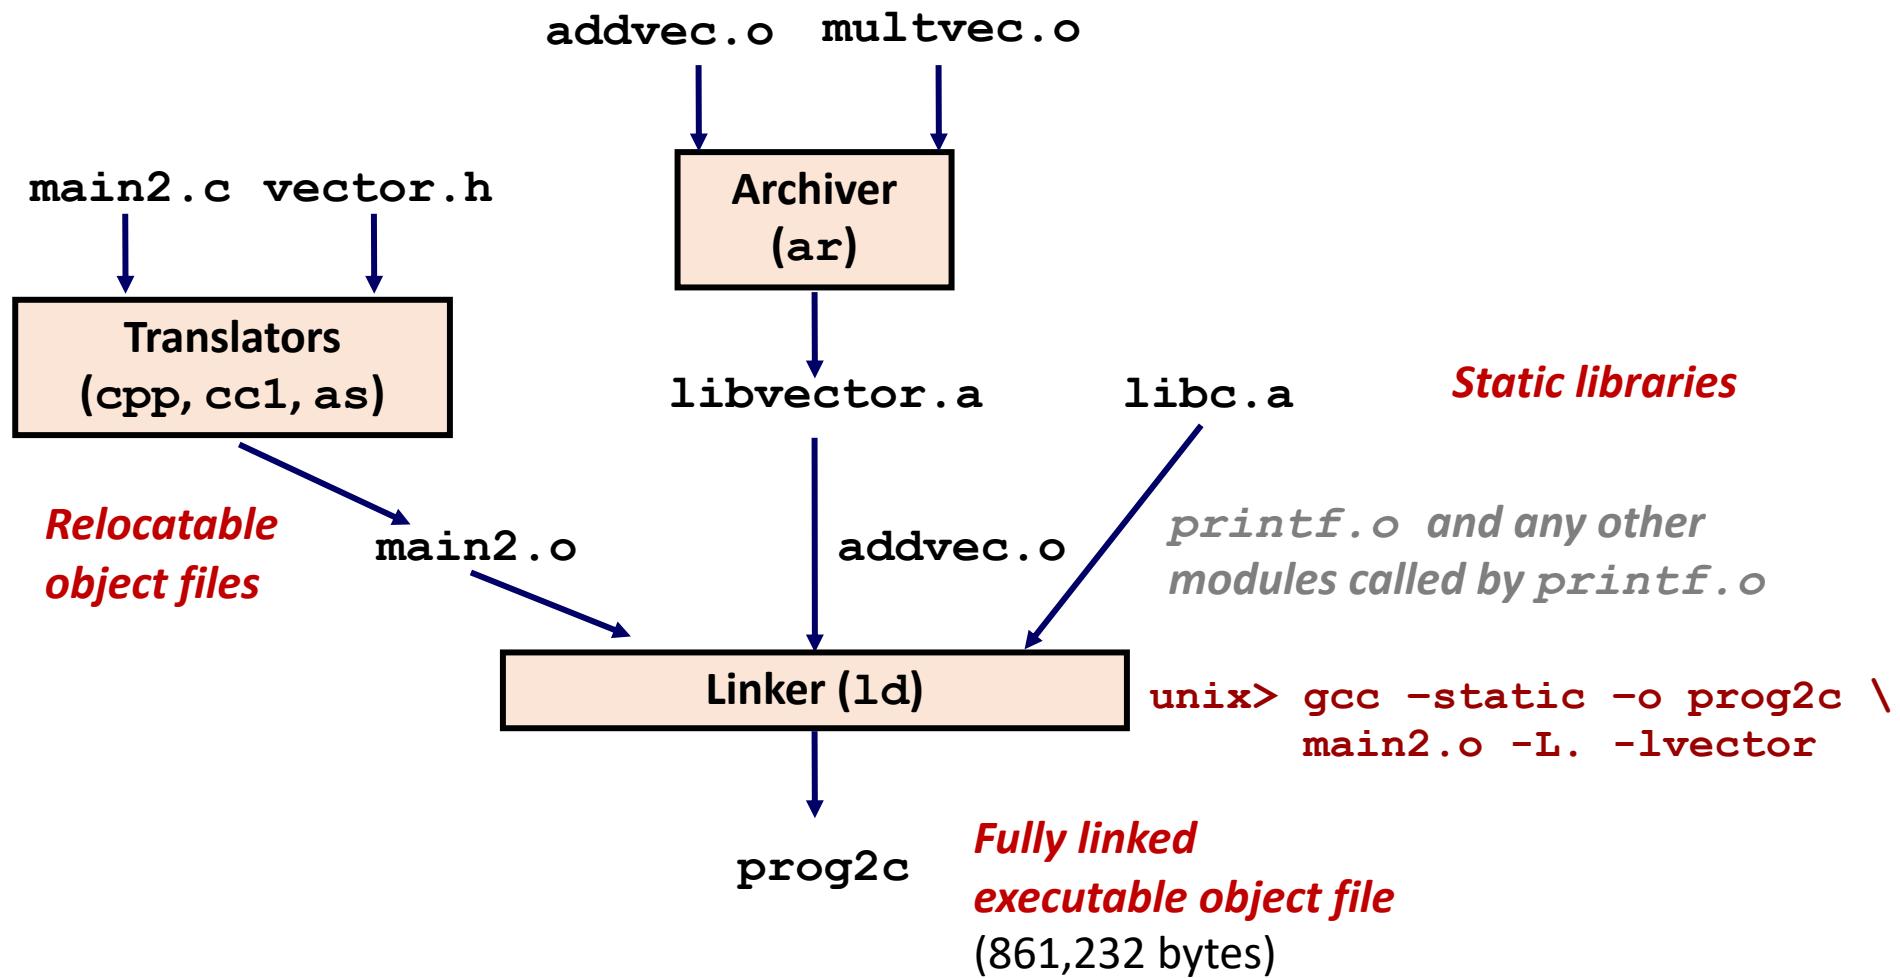
## **libm.a** (the C math library)

- 2 MB archive of 444 object files.
- floating point math (sin, cos, tan, log, exp, sqrt, ...)

```
% ar -t /usr/lib/libc.a | sort
...
fork.o
...
fprintf.o
fpu_control.o
fputc.o
freopen.o
fstab.o
...
```

```
% ar -t /usr/lib/libm.a | sort
...
e_acos.o
e_acosf.o
e_acosh.o
e_acoshf.o
e_acoshl.o
e_asinf.o
e_asinl.o
...
```

# Linking with static libraries



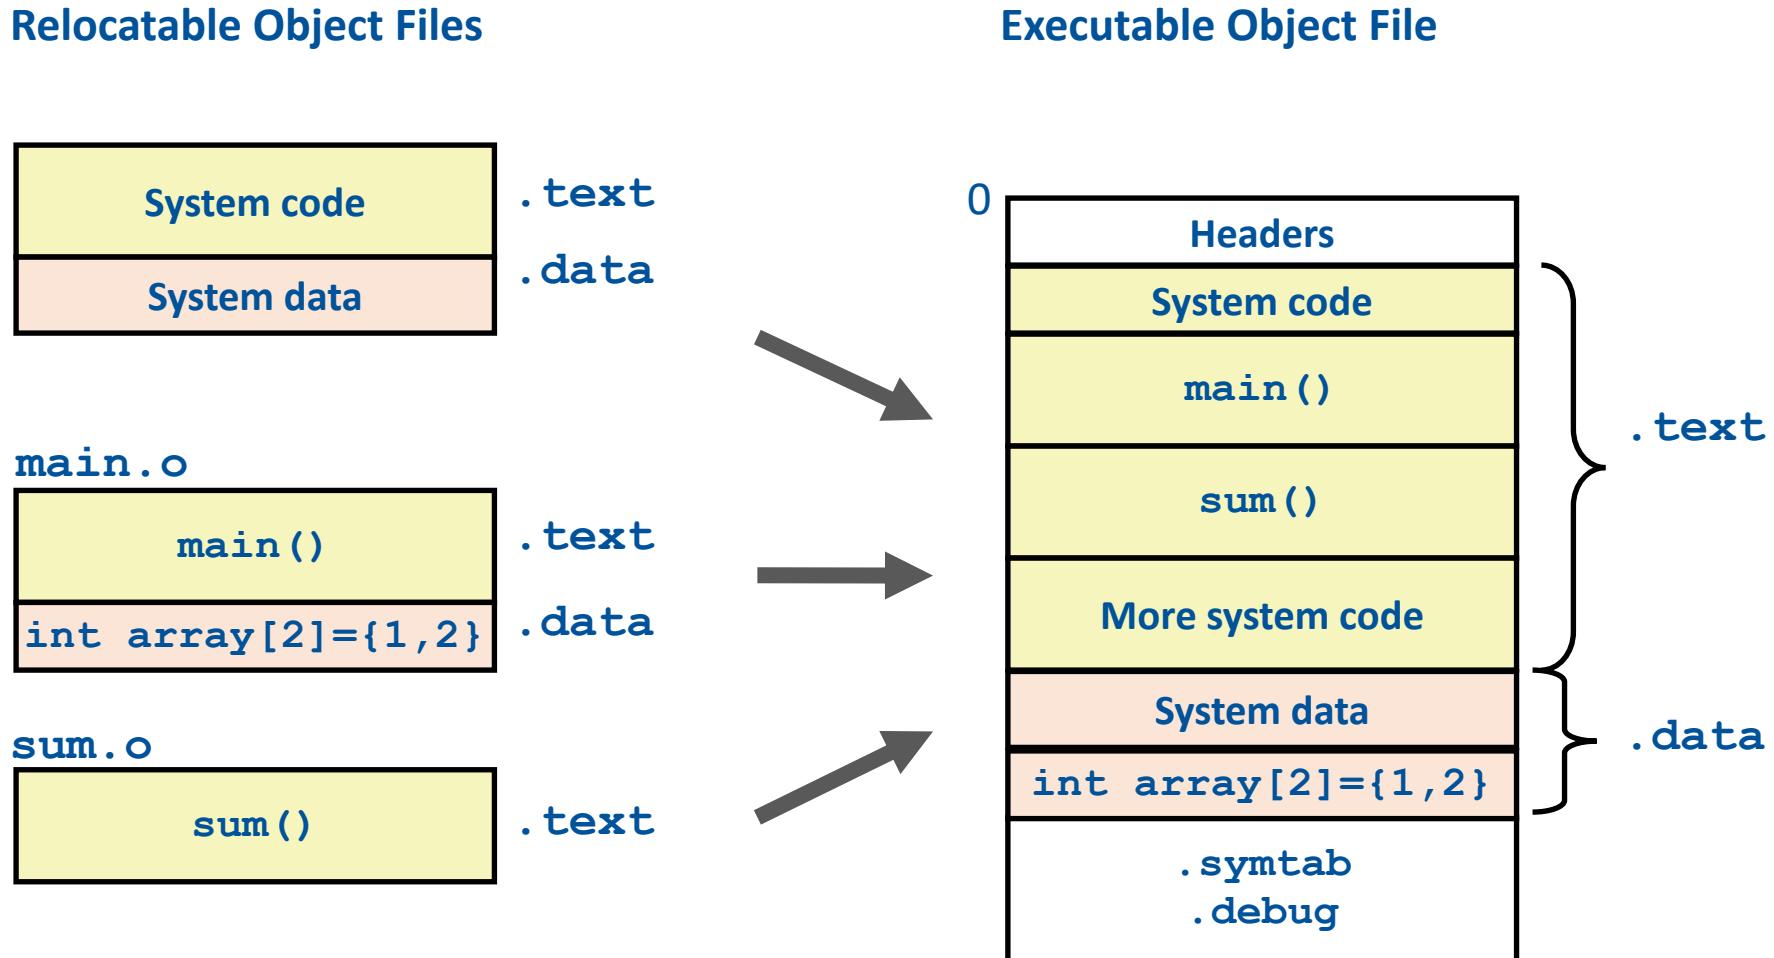
# Static libraries algorithm

- Linker's algorithm for resolving external references:
  - Scan `.o` files and `.a` files in the command line order.
  - During the scan, keep a list of the current unresolved references.
  - As each new `.o` or `.a` file, *obj*, is encountered, try to resolve each unresolved reference in the list against the symbols defined in *obj*.
  - If any entries in the unresolved list at end of scan, then error.

```
unix> gcc -static -o prog2c -L. -lvector main2.o
main2.o: In function `main':
main2.c:(.text+0x19): undefined reference to `addvec'
collect2: error: ld returned 1 exit status
```

- Order matters! -- put libraries at the end of the command line.

# Linker relocation



# Relocation entries

```
int array[2] = {1, 2};

int main(int argc, char**
argv)
{
    int val = sum(array, 2);
    return val;
}
```

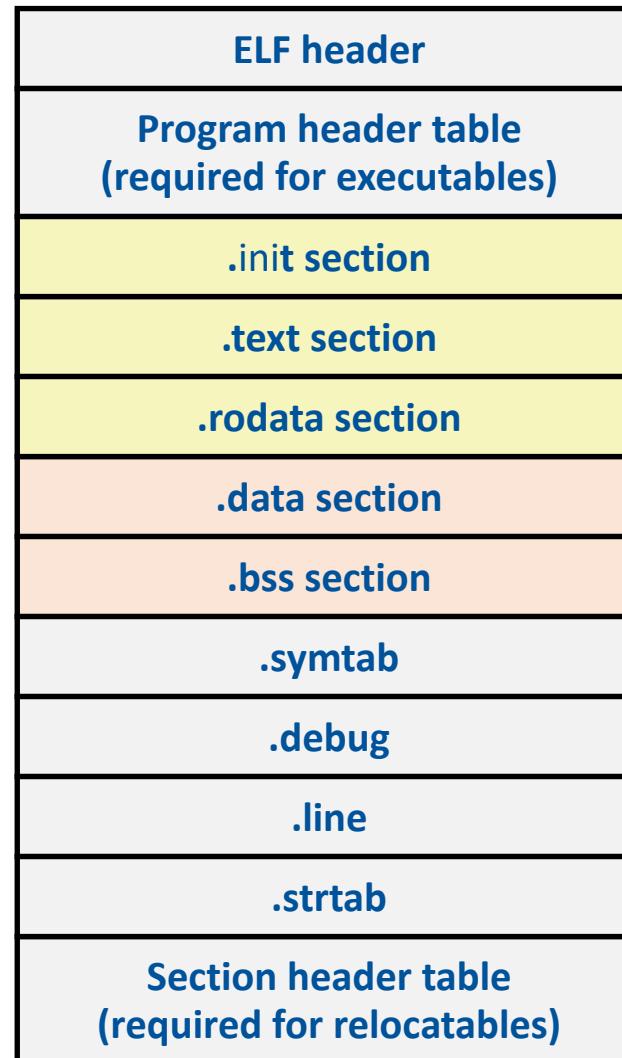
```
0000000000000000 <main>:
 0: 48 83 ec 08          sub    $0x8,%rsp
 4: be 02 00 00 00        mov    $0x2,%esi
 9: bf 00 00 00 00        mov    $0x0,%edi      # %edi = &array
                                     a: R_X86_64_32 array      # Relocation entry

e:  e8 00 00 00 00        callq   13 <main+0x13> # sum()
                                     f: R_X86_64_PC32 sum-0x4      # Relocation entry
13: 48 83 c4 08          add    $0x8,%rsp
17: c3                   retq
```

# Relocated .text section

```
00000000004004d0 <main>:  
4004d0: 48 83 ec 08      sub    $0x8,%rsp  
4004d4: be 02 00 00 00    mov    $0x2,%esi  
4004d9: bf 18 10 60 00    mov    $0x601018,%edi  # %edi = &array  
4004de: e8 05 00 00 00    callq  4004e8 <sum>    # sum()  
4004e3: 48 83 c4 08    add    $0x8,%rsp  
4004e7: c3                retq  
  
00000000004004e8 <sum>:  
4004e8: b8 00 00 00 00    mov    $0x0,%eax  
4004ed: ba 00 00 00 00    mov    $0x0,%edx  
4004f2: eb 09                jmp    4004fd <sum+0x15>  
4004f4: 48 63 ca      movslq %edx,%rcx  
4004f7: 03 04 8f      add    (%rdi,%rcx,4),%eax  
4004fa: 83 c2 01      add    $0x1,%edx  
4004fd: 39 f2                cmp    %esi,%edx  
4004ff: 7c f3                jl    4004f4 <sum+0xc>  
400501: f3 c3      repz    retq
```

# Executable object file



# Next Class

- Dynamic linking
- **Readings:** 7.12 – 7.14